

Import Libraries

In [1]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

In [2]:

```
Churn = pd.read_csv("F:\dataset\Churn_Modelling.csv")
```

In [3]:

```
Churn.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   RowNumber             10000 non-null  int64  
1   CustomerId            10000 non-null  int64  
2   Surname               10000 non-null  object  
3   CreditScore           10000 non-null  int64  
4   Geography             10000 non-null  object  
5   Gender                10000 non-null  object  
6   Age                  10000 non-null  int64  
7   Tenure                10000 non-null  int64  
8   Balance               10000 non-null  float64 
9   NumOfProducts        10000 non-null  int64  
10  HasCrCard             10000 non-null  int64  
11  IsActiveMember        10000 non-null  int64  
12  EstimatedSalary       10000 non-null  float64 
13  Exited                10000 non-null  int64  
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB
```

In [4]:

```
Churn.head()
```

Out[4]:


	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance
0	1	15634602	Hargrave	619	France	Female	42	2	
1	2	15647311	Hill	608	Spain	Female	41	1	8380
2	3	15619304	Onio	502	France	Female	42	8	15966
3	4	15701354	Boni	699	France	Female	39	1	
4	5	15737888	Mitchell	850	Spain	Female	43	2	12551

In [5]:

```
sns.heatmap(Churn.isnull(),yticklabels=False,cbar=False,cmap='viridis')
```

Out[5]:

<matplotlib.axes._subplots.AxesSubplot at 0x265a351a2c8>




In [6]:

```
sns.set_style('whitegrid')  
sns.countplot(x='Exited',data=Churn,palette='RdBu_r')
```

Out[6]:

<matplotlib.axes._subplots.AxesSubplot at 0x265a3d17f88>




In [7]:

```
sns.set_style('whitegrid')
sns.countplot(x='Exited',hue='Gender',data=Churn,palette='RdBu_r')
```

Out[7]:

<matplotlib.axes._subplots.AxesSubplot at 0x265a3dbd288>




In [8]:

```
sns.set_style('whitegrid')
sns.countplot(x='Exited',hue='Geography',data=Churn,palette='rainbow')
```

Out[8]:

<matplotlib.axes._subplots.AxesSubplot at 0x265a40aa9c8>



In [9]:

```
#Churn['Age'].hist(bins=30,color='darkred',alpha=0.7)
```

Data Cleaning

In [10]:

```
to_drop=['RowNumber', 'CustomerId', 'Surname']  
Churn=Churn.drop(to_drop,axis=1)  
Churn.head()
```

Out[10]:


	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	Is
0	619	France	Female	42	2	0.00	1	1	
1	608	Spain	Female	41	1	83807.86	1	0	
2	502	France	Female	42	8	159660.80	3	1	
3	699	France	Female	39	1	0.00	2	0	
4	850	Spain	Female	43	2	125510.82	1	1	

In [11]:

```
plt.figure(figsize=(12, 7))  
sns.boxplot(x='Gender',y='Balance',data=Churn,palette='winter')
```

Out[11]:

<matplotlib.axes._subplots.AxesSubplot at 0x265a411f108>




In [12]:

```
plt.figure(figsize=(12, 7))  
sns.boxplot(x='Geography',y='Balance',data=Churn,palette='winter')
```

Out[12]:

<matplotlib.axes._subplots.AxesSubplot at 0x265a47e74c8>




In [13]:

```
plt.figure(figsize=(12, 7))  
sns.boxplot(x='Gender',y='EstimatedSalary',data=Churn,palette='winter')
```

Out[13]:

<matplotlib.axes._subplots.AxesSubplot at 0x265a411fe88>




In [14]:

```
plt.figure(figsize=(12, 7))  
sns.boxplot(x='Exited',y='CreditScore',data=Churn,palette='winter')
```

Out[14]:

<matplotlib.axes._subplots.AxesSubplot at 0x2659e5d3848>



Converting Categorical Features

In [17]:

```
Churn.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   CreditScore            10000 non-null  int64
1   Geography              10000 non-null  object
2   Gender                 10000 non-null  object
3   Age                   10000 non-null  int64
4   Tenure                 10000 non-null  int64
5   Balance                10000 non-null  float64
6   NumOfProducts          10000 non-null  int64
7   HasCrCard              10000 non-null  int64
8   IsActiveMember         10000 non-null  int64
9   EstimatedSalary        10000 non-null  float64
10  Exited                  10000 non-null  int64
dtypes: float64(2), int64(7), object(2)
memory usage: 859.5+ KB
```

In [18]:

```
sex = pd.get_dummies(Churn['Gender'],drop_first=True)
sex
```

Out[18]:

	Male
0	0
1	0
2	0
3	0
4	0
...	...
9995	1
9996	1
9997	0
9998	1
9999	0

10000 rows × 1 columns

In [19]:

```
Geo = pd.get_dummies(Churn['Geography'],drop_first=True)
Geo
```

Out[19]:

	Germany	Spain
0	0	0
1	0	1
2	0	0
3	0	0
4	0	1
...
9995	0	0
9996	0	0
9997	0	0
9998	1	0
9999	0	0

10000 rows × 2 columns

In [20]:

```
to_drop1=['Gender','Geography']
Churn=Churn.drop(to_drop1,axis=1)
Churn.head()
```

Out[20]:

	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
0	619	42	2	0.00	1	1	1	
1	608	41	1	83807.86	1	0	1	
2	502	42	8	159660.80	3	1	0	
3	699	39	1	0.00	2	0	0	
4	850	43	2	125510.82	1	1	1	

In [21]:

```
Churn = pd.concat([Churn,sex,Geo],axis=1)
```


In [22]:

```
Churn.head()
```

Out[22]:

	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
0	619	42	2	0.00	1	1	1	
1	608	41	1	83807.86	1	0	1	
2	502	42	8	159660.80	3	1	0	
3	699	39	1	0.00	2	0	0	
4	850	43	2	125510.82	1	1	1	



Building a Logistic Regression model

In [23]:

```
from sklearn.model_selection import train_test_split
```

In [24]:

```
X_train, X_test, y_train, y_test = train_test_split(Churn.drop('Exited',axis=1),  
                                                    Churn['Exited'], test_size=0.20,  
                                                    random_state=111)
```

In [25]:

```
print(X_train,y_train)
```

	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	\
27	571	44	9	0.00	2	0	
3847	611	37	6	0.00	2	1	
7461	596	32	4	0.00	2	0	
1356	709	49	4	154344.49	2	1	
4314	638	34	5	133501.36	1	0	
...	
7490	654	35	2	90865.80	1	1	
8873	610	34	0	103108.17	1	0	
7443	634	24	2	87413.19	1	1	
4182	550	52	5	121016.23	1	1	
4820	484	32	3	0.00	2	1	

	IsActiveMember	EstimatedSalary	Male	Germany	Spain
27	0	38433.35	1	0	0
3847	0	110782.88	0	0	0
7461	1	146504.35	1	0	1
1356	1	38794.57	1	0	0
4314	1	155643.04	1	0	0
...
7490	1	86764.46	0	0	0
8873	0	125646.82	1	0	0
7443	0	63340.65	0	0	0
4182	1	41730.37	1	1	0
4820	1	139390.99	0	0	0

```
[8000 rows x 11 columns] 27 0
```

```
3847 0
```

```
7461 0
```

```
1356 0
```

```
4314 0
```

```
..
```

```
7490 0
```

```
8873 0
```

```
7443 0
```

```
4182 1
```

```
4820 0
```

```
Name: Exited, Length: 8000, dtype: int64
```

In [26]:

```
print(X_test,y_test)
```

	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	\
207	618	34	5	134954.53	1	1	
1866	559	70	9	0.00	1	1	
9487	850	32	5	0.00	1	1	
3673	764	24	7	98148.61	1	1	
7178	684	38	5	105069.98	2	1	
...	
3943	649	46	5	0.00	2	1	
4007	648	43	7	139972.18	1	1	
8540	484	40	7	106901.42	2	0	
1906	786	29	4	0.00	2	1	
488	692	30	2	0.00	2	0	

	IsActiveMember	EstimatedSalary	Male	Germany	Spain
207	1	151954.39	1	0	0
1866	1	122996.76	0	0	0
9487	1	3830.59	0	0	1
3673	0	26843.76	1	0	0
7178	1	198355.28	1	0	0
...
3943	1	76946.60	1	0	0
4007	0	143668.58	0	0	0
8540	0	118045.98	1	1	0
1906	0	103372.79	0	0	0
488	1	130486.57	1	0	0

```
[2000 rows x 11 columns] 207      0
```

```
1866      0
```

```
9487      0
```

```
3673      0
```

```
7178      0
```

```
..
```

```
3943      0
```

```
4007      0
```

```
8540      0
```

```
1906      0
```

```
488       0
```

```
Name: Exited, Length: 2000, dtype: int64
```

Training and Predicting

In [27]:

```
from sklearn.linear_model import LogisticRegression
```

In [28]:

```
logmodel = LogisticRegression()
logmodel.fit(X_train,y_train)
```

Out[28]:

```
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                    intercept_scaling=1, l1_ratio=None, max_iter=100,
                    multi_class='auto', n_jobs=None, penalty='l2',
                    random_state=None, solver='lbfgs', tol=0.0001, verbose=
0,
                    warm_start=False)
```

In [29]:

```
predictions = logmodel.predict(X_test)
```

In [30]:

```
predictions
```

Out[30]:

```
array([0, 1, 0, ..., 0, 0, 0], dtype=int64)
```

Let's move on to evaluate our model!

Evaluation

In [31]:

```
from sklearn.metrics import classification_report, confusion_matrix
```

In [32]:

```
print(confusion_matrix(y_test, logmodel.predict(X_test)))
```

```
[[1558  34]
 [ 391  17]]
```


In [33]:

```
print(classification_report(y_test, logmodel.predict(X_test)))
```

	precision	recall	f1-score	support
0	0.80	0.98	0.88	1592
1	0.33	0.04	0.07	408
accuracy			0.79	2000
macro avg	0.57	0.51	0.48	2000
weighted avg	0.70	0.79	0.72	2000

In [34]:

```
from sklearn.metrics import roc_auc_score
from sklearn.metrics import roc_curve
logit_roc_auc = roc_auc_score(y_test, predictions)
fpr, tpr, thresholds = roc_curve(y_test, logmodel.predict_proba(X_test)[:,-1])
plt.figure()
plt.plot(fpr, tpr, label='Logistic Regression (area = %0.2f)' % logit_roc_auc)
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic')
plt.legend(loc="lower right")
#plt.savefig('Log_ROC')
plt.show()
```



In [35]:

```
roc_auc_score(y_test, logmodel.predict(X_test))
```

Out[35]:

0.5101549413735343

In [36]:

```
import statsmodels.api as sm
logit_model=sm.Logit(y_train,X_train)
result=logit_model.fit()
print(result.summary2())
```

Optimization terminated successfully.

Current function value: 0.437496

Iterations 6

Results: Logit

```
=====
Model:                Logit                Pseudo R-squared: 0.134
Dependent Variable:    Exited                AIC:                7021.9371
Date:                  2020-07-06 08:42      BIC:                7098.7962
No. Observations:      8000                Log-Likelihood:     -3500.0
Df Model:              10                  LL-Null:            -4043.1
Df Residuals:          7989                LLR p-value:        4.8508e-227
Converged:              1.0000              Scale:             1.0000
No. Iterations:        6.0000

-----
                Coef.  Std.Err.   z      P>|z|   [0.025  0.975]
-----
CreditScore    -0.0034   0.0002 -16.0292 0.0000 -0.0038 -0.0030
Age             0.0586   0.0026  22.3630 0.0000  0.0534  0.0637
Tenure         -0.0456   0.0101  -4.4962 0.0000 -0.0655 -0.0257
Balance         0.0000   0.0000   1.3242 0.1854 -0.0000  0.0000
NumOfProducts  -0.3786   0.0507  -7.4704 0.0000 -0.4779 -0.2793
HasCrCard      -0.1629   0.0637  -2.5584 0.0105 -0.2877 -0.0381
IsActiveMember -1.0876   0.0633 -17.1750 0.0000 -1.2117 -0.9635
EstimatedSalary -0.0000   0.0000  -1.7459 0.0808 -0.0000  0.0000
Male           -0.6537   0.0597 -10.9529 0.0000 -0.7706 -0.5367
Germany         0.7855   0.0754  10.4220 0.0000  0.6378  0.9333
Spain          -0.0295   0.0774  -0.3813 0.7030 -0.1813  0.1222
=====
```

In [37]:

```
to_drop2=['Balance','EstimatedSalary','Spain']
Churn1=Churn.drop(to_drop2,axis=1)
Churn1.head()
```

Out[37]:

	CreditScore	Age	Tenure	NumOfProducts	HasCrCard	IsActiveMember	Exited	Male	Ge
0	619	42	2	1	1	1	1	0	
1	608	41	1	1	0	1	0	0	
2	502	42	8	3	1	0	1	0	
3	699	39	1	2	0	0	0	0	
4	850	43	2	1	1	1	0	0	

In [38]:

```
X_train, X_test, y_train, y_test = train_test_split(Churn1.drop('Exited',axis=1),
                                                    Churn1['Exited'], test_size=0.20,
                                                    random_state=111)
```

In [39]:

```
import statsmodels.api as sm
logit_model=sm.Logit(y_train,X_train)
result=logit_model.fit()
print(result.summary2())
```

Optimization terminated successfully.

Current function value: 0.437789

Iterations 6

Results: Logit

```
=====
Model:                Logit                Pseudo R-squared: 0.134
Dependent Variable:   Exited                AIC:                7020.6276
Date:                2020-07-06 08:42      BIC:                7076.5252
No. Observations:    8000                Log-Likelihood:    -3502.3
Df Model:            7                  LL-Null:           -4043.1
Df Residuals:        7992              LLR p-value:       2.8493e-229
Converged:           1.0000              Scale:            1.0000
No. Iterations:      6.0000

-----
              Coef.  Std.Err.   z      P>|z|    [0.025  0.975]
-----
CreditScore   -0.0034   0.0002 -17.2715 0.0000 -0.0038 -0.0030
Age            0.0584   0.0026  22.4868 0.0000  0.0533  0.0634
Tenure        -0.0457   0.0101  -4.5223 0.0000 -0.0656 -0.0259
NumOfProducts -0.4005   0.0493  -8.1207 0.0000 -0.4972 -0.3038
HasCrCard     -0.1645   0.0636  -2.5866 0.0097 -0.2892 -0.0399
IsActiveMember -1.0844   0.0633 -17.1393 0.0000 -1.2084 -0.9604
Male          -0.6510   0.0595 -10.9325 0.0000 -0.7677 -0.5343
Germany       0.8332   0.0641  12.9964 0.0000  0.7075  0.9589
=====
```

In [40]:

```
logmodel = LogisticRegression()
logmodel.fit(X_train,y_train)
predictions1 = logmodel.predict(X_test)
predictions1
```

C:\Users\Pratima Dhar\anaconda3\lib\site-packages\sklearn\linear_model_logistic.py:940: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)

Out[40]:

```
array([0, 1, 0, ..., 0, 0, 0], dtype=int64)
```

In [41]:

```
from sklearn.metrics import classification_report, confusion_matrix
```

In [42]:

```
print(confusion_matrix(y_test, logmodel.predict(X_test)))
```

```
[[1534  58]
 [ 325  83]]
```

In [43]:

```
print(classification_report(y_test, logmodel.predict(X_test)))
```

	precision	recall	f1-score	support
0	0.83	0.96	0.89	1592
1	0.59	0.20	0.30	408
accuracy			0.81	2000
macro avg	0.71	0.58	0.60	2000
weighted avg	0.78	0.81	0.77	2000

In [44]:

```
from sklearn.metrics import roc_auc_score
from sklearn.metrics import roc_curve
logit_roc_auc = roc_auc_score(y_test, predictions1)
fpr, tpr, thresholds = roc_curve(y_test, logmodel.predict_proba(X_test)[: ,1])
plt.figure()
plt.plot(fpr, tpr, label='Logistic Regression (area = %0.2f)' % logit_roc_auc)
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic')
plt.legend(loc="lower right")
#plt.savefig('Log_ROC')
plt.show()
```

