

Import libraries

In [2]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

In [3]:

```
iris = pd.read_csv("F:\dataset\iris.csv")
```

In [4]:

```
iris.head()
```

Out[4]:

	SL	SW	PL	PW	Classification
0	5.1	3.5	1.4	0.2	Iris-setosa
1	4.9	3.0	1.4	0.2	Iris-setosa
2	4.7	3.2	1.3	0.2	Iris-setosa
3	4.6	3.1	1.5	0.2	Iris-setosa
4	5.0	3.6	1.4	0.2	Iris-setosa

In [5]:

```
iris.shape
```

Out[5]:

```
(150, 5)
```

In [6]:

```
iris.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  -
0    SL              150 non-null   float64
1    SW              150 non-null   float64
2    PL              150 non-null   float64
3    PW              150 non-null   float64
4    Classification  150 non-null   object
dtypes: float64(4), object(1)
memory usage: 6.0+ KB
```

In [7]:

```
iris.isnull()
```

Out[7]:

	SL	SW	PL	PW	Classification
0	False	False	False	False	False
1	False	False	False	False	False
2	False	False	False	False	False
3	False	False	False	False	False
4	False	False	False	False	False
...
145	False	False	False	False	False
146	False	False	False	False	False
147	False	False	False	False	False
148	False	False	False	False	False
149	False	False	False	False	False


150 rows × 5 columns

In [8]:

```
plt.figure(figsize=(12, 7))
sns.boxplot(x='Classification',y='SL',data=iris,palette='winter')
```

Out[8]:

<matplotlib.axes._subplots.AxesSubplot at 0x295b69b45c8>




In [9]:

```
plt.figure(figsize=(12, 7))  
sns.boxplot(x='Classification',y='PL',data=iris,palette='winter')
```

Out[9]:

<matplotlib.axes._subplots.AxesSubplot at 0x295b81e5bc8>




In [10]:

```
plt.figure(figsize=(12, 7))  
sns.boxplot(x='Classification',y='SW',data=iris,palette='winter')
```

Out[10]:

<matplotlib.axes._subplots.AxesSubplot at 0x295b8243288>




In [11]:

```
plt.figure(figsize=(12, 7))  
sns.boxplot(x='Classification', y='PW', data=iris, palette='winter')
```

Out[11]:

<matplotlib.axes._subplots.AxesSubplot at 0x295b85e0bc8>



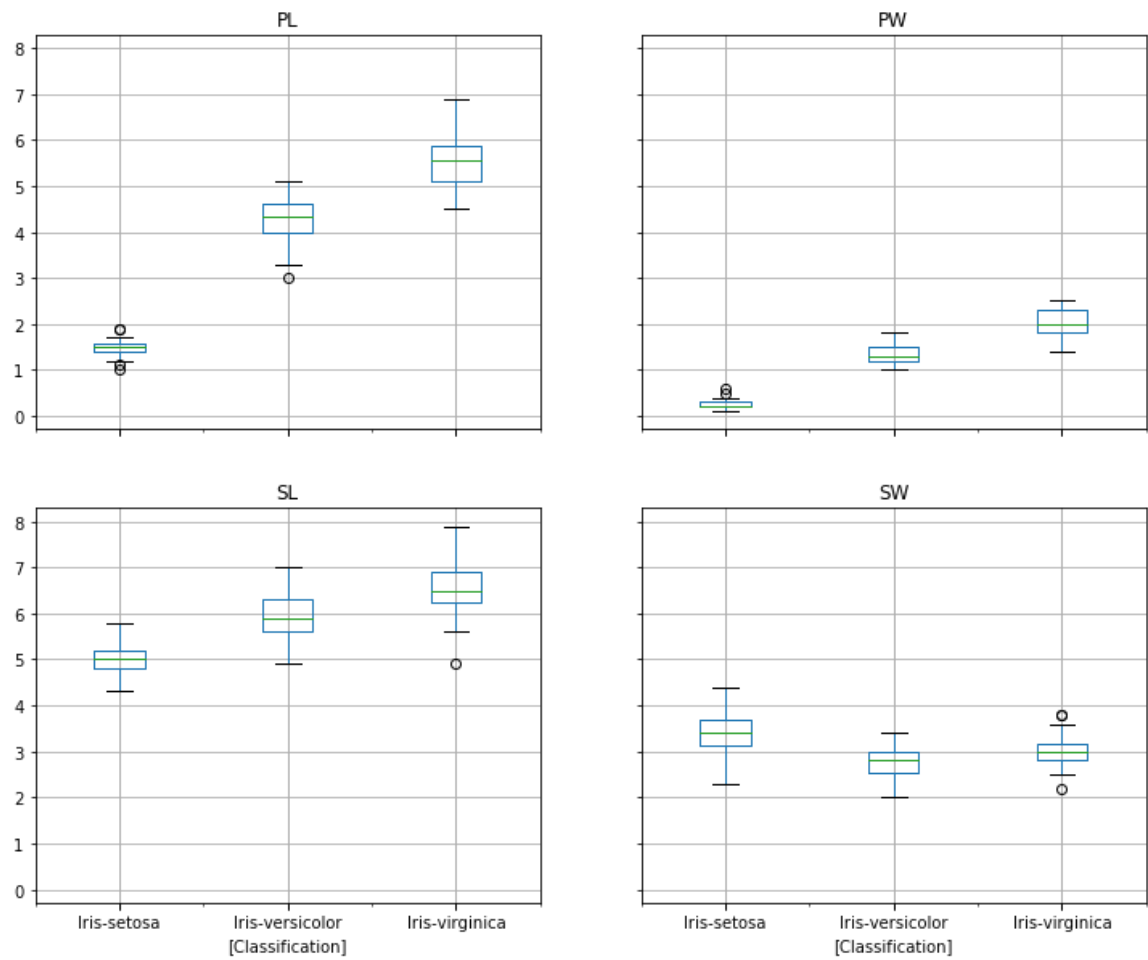
In [12]:

```
iris.boxplot(by="Classification", figsize=(12, 10))
```

Out[12]:

```
array([[<matplotlib.axes._subplots.AxesSubplot object at 0x00000295B86C118
8>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x00000295B86EFCC
8>],
      [<matplotlib.axes._subplots.AxesSubplot object at 0x00000295B86C14C
8>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x00000295B876128
8>]],
      dtype=object)
```

Boxplot grouped by Classification



In [13]:

```
sns.pairplot(iris, hue="Classification", palette="husl", size=3, diag_kind="kde")
```

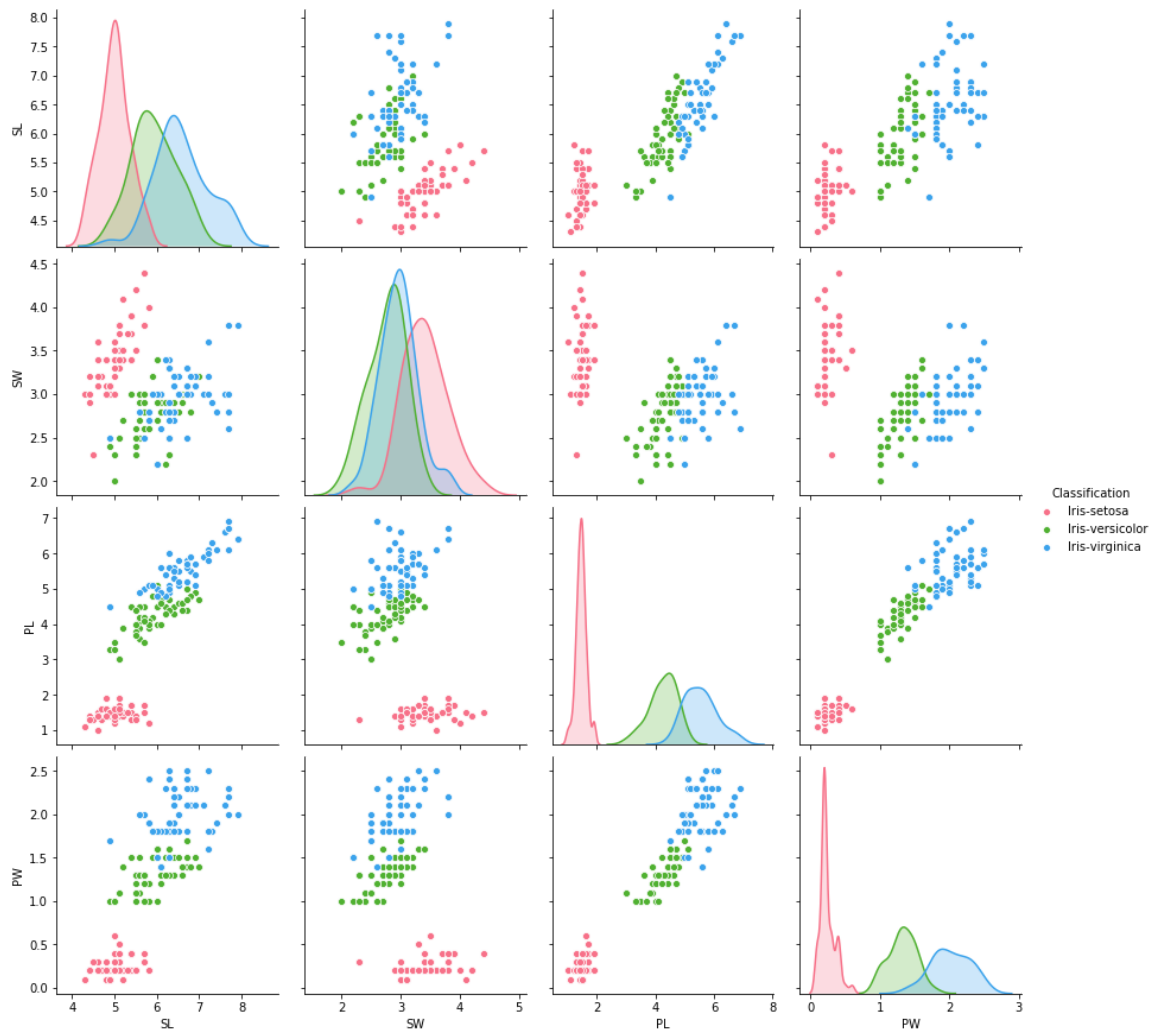
C:\Users\Pratima Dhar\anaconda3\lib\site-packages\seaborn\axisgrid.py:207

9: UserWarning: The `size` parameter has been renamed to `height`; please update your code.

```
warnings.warn(msg, UserWarning)
```

Out[13]:

```
<seaborn.axisgrid.PairGrid at 0x295b8c10948>
```



In [14]:

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(iris.drop('Classification',axis=1),
                                                    iris['Classification'], test_size=
0.20,
                                                    random_state=10)
```

In [15]:

```
print(X_train,y_train)
```

	SL	SW	PL	PW
58	6.6	2.9	4.6	1.3
97	6.2	2.9	4.3	1.3
129	7.2	3.0	5.8	1.6
114	5.8	2.8	5.1	2.4
146	6.3	2.5	5.0	1.9
..
113	5.7	2.5	5.0	2.0
64	5.6	2.9	3.6	1.3
15	5.7	4.4	1.5	0.4
125	7.2	3.2	6.0	1.8
9	4.9	3.1	1.5	0.1

```
[120 rows x 4 columns] 58      Iris-versicolor
97      Iris-versicolor
129     Iris-virginica
114     Iris-virginica
146     Iris-virginica
...
113     Iris-virginica
64      Iris-versicolor
15      Iris-setosa
125     Iris-virginica
9       Iris-setosa
Name: Classification, Length: 120, dtype: object
```


In [16]:

```
print(X_test,y_test)
```

	SL	SW	PL	PW	
87	6.3	2.3	4.4	1.3	
111	6.4	2.7	5.3	1.9	
10	5.4	3.7	1.5	0.2	
91	6.1	3.0	4.6	1.4	
49	5.0	3.3	1.4	0.2	
60	5.0	2.0	3.5	1.0	
72	6.3	2.5	4.9	1.5	
67	5.8	2.7	4.1	1.0	
39	5.1	3.4	1.5	0.2	
55	5.7	2.8	4.5	1.3	
66	5.6	3.0	4.5	1.5	
142	5.8	2.7	5.1	1.9	
53	5.5	2.3	4.0	1.3	
1	4.9	3.0	1.4	0.2	
19	5.1	3.8	1.5	0.3	
112	6.8	3.0	5.5	2.1	
85	6.0	3.4	4.5	1.6	
38	4.4	3.0	1.3	0.2	
21	5.1	3.7	1.5	0.4	
35	5.0	3.2	1.2	0.2	
102	7.1	3.0	5.9	2.1	
132	6.4	2.8	5.6	2.2	
126	6.2	2.8	4.8	1.8	
24	4.8	3.4	1.9	0.2	
61	5.9	3.0	4.2	1.5	
2	4.7	3.2	1.3	0.2	
95	5.7	3.0	4.2	1.2	
90	5.5	2.6	4.4	1.2	
76	6.8	2.8	4.8	1.4	
117	7.7	3.8	6.7	2.2	87 Iris-versicolor
111	Iris-virginica				
10	Iris-setosa				
91	Iris-versicolor				
49	Iris-setosa				
60	Iris-versicolor				
72	Iris-versicolor				
67	Iris-versicolor				
39	Iris-setosa				
55	Iris-versicolor				
66	Iris-versicolor				
142	Iris-virginica				
53	Iris-versicolor				
1	Iris-setosa				
19	Iris-setosa				
112	Iris-virginica				
85	Iris-versicolor				
38	Iris-setosa				
21	Iris-setosa				
35	Iris-setosa				
102	Iris-virginica				
132	Iris-virginica				
126	Iris-virginica				
24	Iris-setosa				
61	Iris-versicolor				
2	Iris-setosa				
95	Iris-versicolor				
90	Iris-versicolor				
76	Iris-versicolor				
117	Iris-virginica				

Name: Classification, dtype: object

In [17]:

```
from sklearn.neighbors import KNeighborsClassifier
knn=KNeighborsClassifier(n_neighbors=4)
knn.fit(X_train,y_train)
```

Out[17]:

```
KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
                     metric_params=None, n_jobs=None, n_neighbors=4, p=2,
                     weights='uniform')
```

In [18]:

```
pred1=knn.predict(X_test)
```

In [19]:

```
print(pred1)
```

```
['Iris-versicolor' 'Iris-virginica' 'Iris-setosa' 'Iris-versicolor'
'Iris-setosa' 'Iris-versicolor' 'Iris-virginica' 'Iris-versicolor'
'Iris-setosa' 'Iris-versicolor' 'Iris-versicolor' 'Iris-virginica'
'Iris-versicolor' 'Iris-setosa' 'Iris-setosa' 'Iris-virginica'
'Iris-versicolor' 'Iris-setosa' 'Iris-setosa' 'Iris-setosa'
'Iris-virginica' 'Iris-virginica' 'Iris-virginica' 'Iris-setosa'
'Iris-versicolor' 'Iris-setosa' 'Iris-versicolor' 'Iris-versicolor'
'Iris-versicolor' 'Iris-virginica']
```

In [20]:

```
from sklearn.metrics import classification_report,confusion_matrix
print(confusion_matrix(y_test,knn.predict(X_test)))
```

```
[[10  0  0]
 [ 0 12  1]
 [ 0  0  7]]
```

In [21]:

```
print(classification_report(y_test,knn.predict(X_test)))
```

	precision	recall	f1-score	support
Iris-setosa	1.00	1.00	1.00	10
Iris-versicolor	1.00	0.92	0.96	13
Iris-virginica	0.88	1.00	0.93	7
accuracy			0.97	30
macro avg	0.96	0.97	0.96	30
weighted avg	0.97	0.97	0.97	30

In [22]:

```
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import cross_val_score
from sklearn.metrics import accuracy_score
from sklearn.metrics import classification_report
from sklearn.metrics import recall_score , precision_score , roc_auc_score , roc_curve
from sklearn.metrics import confusion_matrix
```

In [23]:

```
cv_scores = []
neighbors = list(np.arange(3,50,2))
for n in neighbors:
    knn = KNeighborsClassifier(n_neighbors = n,algorithm = 'brute')

    cross_val = cross_val_score(knn,X_train,y_train,cv = 5 , scoring = 'accuracy')
    cv_scores.append(cross_val.mean())

error = [1-x for x in cv_scores]
optimal_n = neighbors[ error.index(min(error)) ]
knn_optimal = KNeighborsClassifier(n_neighbors = optimal_n,algorithm = 'brute')
knn_optimal.fit(X_train,y_train)
pred = knn_optimal.predict(X_test)
acc = accuracy_score(y_test,pred)*100
print("The accuracy for optimal k = {0} using brute is {1}".format(optimal_n,acc))
```

The accuracy for optimal k = 7 using brute is 96.66666666666667

In [24]:


```
from sklearn.neighbors import KNeighborsClassifier
from sklearn import metrics
k_range=range(1,25)
scores={}
scores_list=[]
for k in k_range:
    knn=KNeighborsClassifier(n_neighbors=k)
    knn.fit(X_train,y_train)
    pred2=knn.predict(X_test)
    scores[k]=metrics.accuracy_score(y_test,pred2)
    scores_list.append(metrics.accuracy_score(y_test,pred2))
```

In [25]:

```
%matplotlib inline
import matplotlib.pyplot as plt
plt.plot(k_range,scores_list)
plt.xlabel('Value of K for KNN')
plt.ylabel('Testing accuracy')
```

Out[25]:

Text(0, 0.5, 'Testing accuracy')



In [26]:

```
from sklearn.neighbors import KNeighborsClassifier
knn=KNeighborsClassifier(n_neighbors=13)
knn.fit(X_train,y_train)
```

Out[26]:

```
KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
                    metric_params=None, n_jobs=None, n_neighbors=13, p=2,
                    weights='uniform')
```

In [27]:

```
pred3=knn.predict(X_test)
```

In [28]:

```
print(pred3)
```

```
['Iris-versicolor' 'Iris-virginica' 'Iris-setosa' 'Iris-versicolor'
'Iris-setosa' 'Iris-versicolor' 'Iris-versicolor' 'Iris-versicolor'
'Iris-setosa' 'Iris-versicolor' 'Iris-versicolor' 'Iris-virginica'
'Iris-versicolor' 'Iris-setosa' 'Iris-setosa' 'Iris-virginica'
'Iris-versicolor' 'Iris-setosa' 'Iris-setosa' 'Iris-setosa'
'Iris-virginica' 'Iris-virginica' 'Iris-virginica' 'Iris-setosa'
'Iris-versicolor' 'Iris-setosa' 'Iris-versicolor' 'Iris-versicolor'
'Iris-versicolor' 'Iris-virginica']
```

In [29]:

```
from sklearn.metrics import classification_report, confusion_matrix
print(confusion_matrix(y_test, knn.predict(X_test)))
```

```
[[10  0  0]
 [ 0 13  0]
 [ 0  0  7]]
```

In [30]:

```
print(classification_report(y_test, knn.predict(X_test)))
```

	precision	recall	f1-score	support
Iris-setosa	1.00	1.00	1.00	10
Iris-versicolor	1.00	1.00	1.00	13
Iris-virginica	1.00	1.00	1.00	7
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30

In [31]:

```
to_drop=['SL', 'SW']
iris1=iris.drop(to_drop,axis=1)
iris1.head()
X_train, X_test, y_train, y_test = train_test_split(iris1.drop('Classification',axis=1),
                                                    iris1['Classification'], test_size=
0.20,
                                                    random_state=21)
```

In [32]:

```
print(X_train,y_train)
```

	PL	PW
41	1.3	0.3
131	6.4	2.0
70	4.8	1.8
46	1.6	0.2
126	4.8	1.8
..
120	5.7	2.3
112	5.5	2.1
48	1.5	0.2
4	1.4	0.2
56	4.7	1.6

```
[120 rows x 2 columns] 41      Iris-setosa
131      Iris-virginica
70      Iris-versicolor
46      Iris-setosa
126      Iris-virginica
...
120      Iris-virginica
112      Iris-virginica
48      Iris-setosa
4      Iris-setosa
56      Iris-versicolor
Name: Classification, Length: 120, dtype: object
```

In [33]:

```
print(X_test,y_test)
```


	PL	PW	
92	4.0	1.2	
44	1.9	0.4	
7	1.5	0.2	
21	1.5	0.4	
95	4.2	1.2	
75	4.4	1.4	
20	1.7	0.2	
121	4.9	2.0	
26	1.6	0.4	
19	1.5	0.3	
81	3.7	1.0	
88	4.1	1.3	
143	5.9	2.3	
117	6.7	2.2	
23	1.7	0.5	
77	5.0	1.7	
138	4.8	1.8	
73	4.7	1.2	
14	1.2	0.2	
142	5.1	1.9	
123	4.9	1.8	
62	4.0	1.0	
83	5.1	1.6	
74	4.3	1.3	
42	1.3	0.2	
60	3.5	1.0	
40	1.3	0.3	
45	1.4	0.3	
87	4.4	1.3	
124	5.7	2.1	92 Iris-versicolor
44			Iris-setosa
7			Iris-setosa
21			Iris-setosa
95			Iris-versicolor
75			Iris-versicolor
20			Iris-setosa
121			Iris-virginica
26			Iris-setosa
19			Iris-setosa
81			Iris-versicolor
88			Iris-versicolor
143			Iris-virginica
117			Iris-virginica
23			Iris-setosa
77			Iris-versicolor
138			Iris-virginica
73			Iris-versicolor
14			Iris-setosa
142			Iris-virginica
123			Iris-virginica
62			Iris-versicolor
83			Iris-versicolor
74			Iris-versicolor
42			Iris-setosa
60			Iris-versicolor
40			Iris-setosa
45			Iris-setosa
87			Iris-versicolor
124			Iris-virginica

Name: Classification, dtype: object

In [34]:

```
from sklearn.neighbors import KNeighborsClassifier
knn2=KNeighborsClassifier(n_neighbors=5)
knn2.fit(X_train,y_train)
```

Out[34]:

```
KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
                    metric_params=None, n_jobs=None, n_neighbors=5, p=2,
                    weights='uniform')
```

In [35]:

```
pred4=knn2.predict(X_test)
```

In [36]:

```
print(pred4)
```

```
['Iris-versicolor' 'Iris-setosa' 'Iris-setosa' 'Iris-setosa'
'Iris-versicolor' 'Iris-versicolor' 'Iris-setosa' 'Iris-virginica'
'Iris-setosa' 'Iris-setosa' 'Iris-versicolor' 'Iris-versicolor'
'Iris-virginica' 'Iris-virginica' 'Iris-setosa' 'Iris-virginica'
'Iris-virginica' 'Iris-versicolor' 'Iris-setosa' 'Iris-virginica'
'Iris-virginica' 'Iris-versicolor' 'Iris-virginica' 'Iris-versicolor'
'Iris-setosa' 'Iris-versicolor' 'Iris-setosa' 'Iris-setosa'
'Iris-versicolor' 'Iris-virginica']
```

In [37]:

```
from sklearn.metrics import classification_report,confusion_matrix
print(confusion_matrix(y_test,knn2.predict(X_test)))
```

```
[[11  0  0]
 [ 0 10  2]
 [ 0  0  7]]
```

In [38]:

```
print(classification_report(y_test,knn2.predict(X_test)))
```

	precision	recall	f1-score	support
Iris-setosa	1.00	1.00	1.00	11
Iris-versicolor	1.00	0.83	0.91	12
Iris-virginica	0.78	1.00	0.88	7
accuracy			0.93	30
macro avg	0.93	0.94	0.93	30
weighted avg	0.95	0.93	0.93	30

In [39]:


```
from sklearn.neighbors import KNeighborsClassifier
from sklearn import metrics
k_range=range(1,25)
scores={}
scores_list=[]
for k in k_range:
    knn=KNeighborsClassifier(n_neighbors=k)
    knn.fit(X_train,y_train)
    pred5=knn.predict(X_test)
    scores[k]=metrics.accuracy_score(y_test,pred5)
    scores_list.append(metrics.accuracy_score(y_test,pred5))
```

In [40]:

```
%matplotlib inline
import matplotlib.pyplot as plt
plt.plot(k_range,scores_list)
plt.xlabel('Value of K for KNN')
plt.ylabel('Testing accuracy')
```

Out[40]:

Text(0, 0.5, 'Testing accuracy')



In [41]:

```
from sklearn.neighbors import KNeighborsClassifier
knn3=KNeighborsClassifier(n_neighbors=15)
knn3.fit(X_train,y_train)
```

Out[41]:

```
KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
                    metric_params=None, n_jobs=None, n_neighbors=15, p=2,
                    weights='uniform')
```

In [42]:

```
pred6=knn3.predict(X_test)
```

In [43]:

```
print(pred6)
```

```
['Iris-versicolor' 'Iris-setosa' 'Iris-setosa' 'Iris-setosa'
'Iris-versicolor' 'Iris-versicolor' 'Iris-setosa' 'Iris-virginica'
'Iris-setosa' 'Iris-setosa' 'Iris-versicolor' 'Iris-versicolor'
'Iris-virginica' 'Iris-virginica' 'Iris-setosa' 'Iris-virginica'
'Iris-virginica' 'Iris-versicolor' 'Iris-setosa' 'Iris-virginica'
'Iris-virginica' 'Iris-versicolor' 'Iris-virginica' 'Iris-versicolor'
'Iris-setosa' 'Iris-versicolor' 'Iris-setosa' 'Iris-setosa'
'Iris-versicolor' 'Iris-virginica']
```

In [44]:

```
from sklearn.metrics import classification_report, confusion_matrix
print(confusion_matrix(y_test, knn3.predict(X_test)))
```

```
[[11  0  0]
 [ 0 10  2]
 [ 0  0  7]]
```

In [45]:

```
print(classification_report(y_test, knn3.predict(X_test)))
```

	precision	recall	f1-score	support
Iris-setosa	1.00	1.00	1.00	11
Iris-versicolor	1.00	0.83	0.91	12
Iris-virginica	0.78	1.00	0.88	7
accuracy			0.93	30
macro avg	0.93	0.94	0.93	30
weighted avg	0.95	0.93	0.93	30